

Introduction to Machine Learning.

Statistical foundations

SUMMARY

Trustworthy AI (TAI)	2
1. Machine Learning	3
2. Types of learning	5
3. The learning problem	6
4. Ensemble learning.....	8
5. Probability, statistics and information theory	10

Artificial Intelligence (AI)

- Symbolic AI
- Computational Intelligence (CI)

Data mining (DM)

Machine Learning (ML)

- Deep Learning (DL)

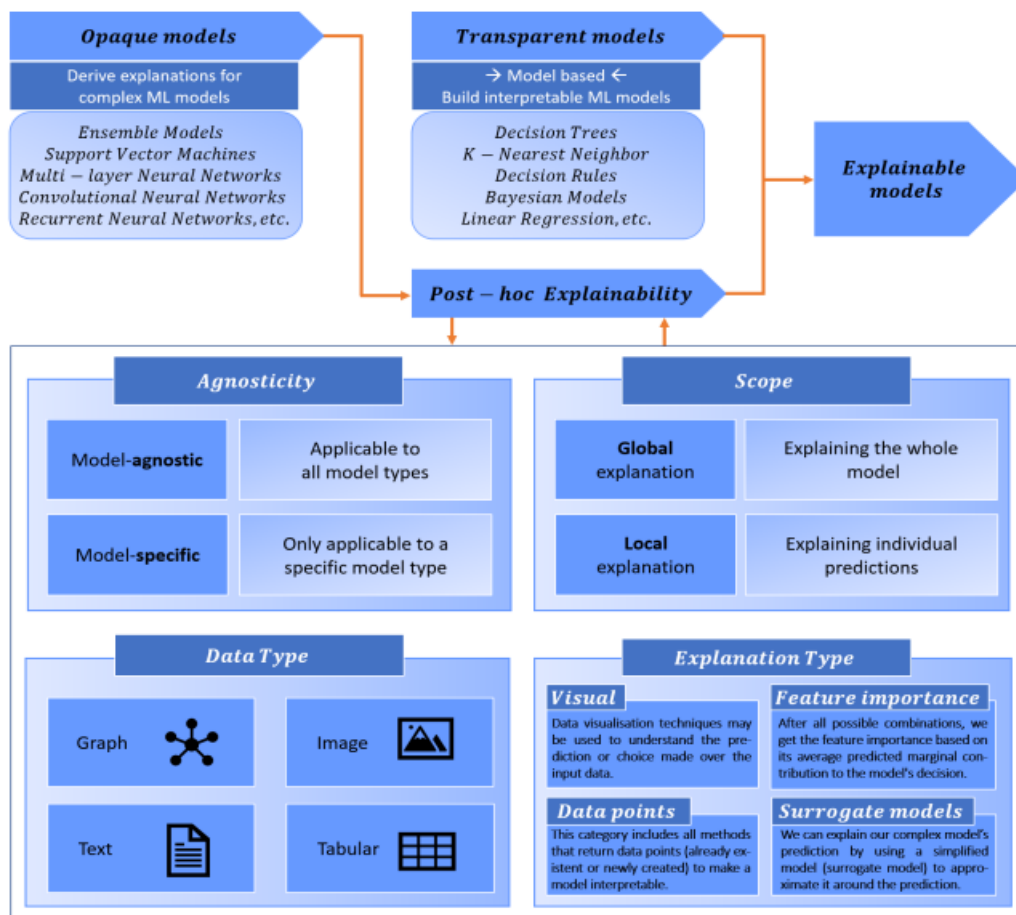
ML/DL

- rapid and significant advancements
 - natural language processing, computer vision, etc
- “hot” topics
 - **Large Language Models** (LLMs)
 - neural network-based models
 - inspired by the human brain
 - trained DL model that understands and generates text in a human-like manner
 - GPT-like models
 - **Generative Pre-trained Transformer**
 - Natural language processing (NLP), genomics, etc
 - [Trustworthiness?](#)
 - **Generative AI/ML**
 - [Machine Unlearning](#)
 - Inspiration – **human brain**
 - **neuroscience-based adaptive unlearning process**
 - sensitive data in training
 - the process of selectively removing specific training data points and their influence on an already trained model, making the updated model behave the same as a model that was never trained on that data
 - [Representation/feature learning](#)
 - Feature fusion
 - Computer vision tasks, bioinformatics, etc
 - [Contrastive learning](#)

- training a model to differentiate between similar and dissimilar pairs of data instances
 - maximizing the similarity within the same class and minimizing the similarity between different classes.

AI

- interpretability of ML/DL models
 - black-box vs white-box models
- [Explainable AI \(XAI\)](#)
 - understand and interpret the predictions of a ML model
- [interpretability methods](#) and evaluation of ML interpretability
 - Gradients, DeepLIFT, Class Activation Maps (CAMs), etc
- methods for ML interpretability can be classified according to various criteria



General ontology of XAI ([link](#))

- [LIME](#) (Local Interpretable Model-agnostic Explanations)
 - algorithm that can explain individual predictions of any classifier or regressor
- [SHAP](#) (SHapley Additive exPlanations)
 - a game theoretic approach to explain the output of any machine learning model

[Trustworthy AI \(TAI\)](#)

- not only **metrics**
- principle, tools

- [ethics guidelines for TAI](#)
- **other aspects**
 - robustness, algorithmic fairness, explainability, and transparency
 - [robustness](#)
 - requirement of TAI
 - refers to the degree that a model's performance changes when using new data versus training data
 - ideally, performance should not deviate significantly
 - the ability of a model to maintain its performance when faced with uncertainties or adversarial conditions
 - noisy data, distribution shifts, and adversarial attacks
 - how to test robustness?
 - cross-validation
 - adding noisy data
 - ...

1. Machine Learning

Learning problems represent an important research direction in AI, *machine learning*.

Machine Learning (ML)= the study of system models, that based on a set of data (training data) **improve** their **performance** (on a specific **task**) by **experiences** and by learning some specific domain knowledge.

Three main directions for ML

- **Data mining** – extract knowledge from data (use historical data to improve decisions, predictions, etc).
- **Software applications** we cannot program by hand
 - autonomous driving
 - speech recognition
 - handwriting recognition
 - game playing
 - etc.
- **Self customizing programs** – programs that adapt to changing conditions
 - Learn the users' interests
- The attempt of modeling the human reasoning leads to the notion of *intelligent reasoning*.
 - Most of the AI systems are *deductive* ones, capable to draw conclusions (make inferences) based on their initial (or supplied) knowledge, without having the capability to generate new knowledge
 - In situations in which a system has incomplete information (knowledge) about its environment, **LEARNING** is the only way the system could get the needed knowledge.
 - The **learning** assures the **autonomy** of a system (the ability to decide which action to perform without external intervention).
- **Learning to represent a function** (target function)

- ML algorithms are also called *function approximators*
- $f : Inputs \rightarrow Outputs - target$ function
- An input $x \in Inputs$ is characterized by a set of **features** (relevant characteristics of the input)
- Learning goal: to find $h \approx f$, h is called **hypothesis**
- Learning = searching for the *hypothesis* that best fits the data

Feature engineering = create features for ML

- **Manual** feature engineering
 - manually create/select relevant features characterizing the data
 - traditional approach – *classical* (traditional, conventional) **ML models**
- **Automated** feature engineering
 - automatically extracting useful and meaningful features from data
 - using unsupervised learning models (e.g. Principal Component Analysis – PCA, autoencoders, aso)
 - using *deep learning* (DL) models
- **Research** topics
 - unsupervised feature/representation learning
 - feature fusion
 - a.s.o

Deep learning (DL)

- is a subfield of Machine Learning
 - mainly based on ANNs
- hierarchical feature learning
- used for **learning data representations**
 - based on representation learning
 - automatically extract features from raw data
 - opposed to task specific algorithms
 - learn from representative examples

ML applications:

- Medicine, bioinformatics, psychology.
- Music composition, archaeology.
- Software engineering (*Search-based software engineering*)
- Computational photography.
- Computer Vision.
- Natural language processing.
- Meteorology.
- Educational data mining.
- aso

Relevant disciplines

- Artificial intelligence
 - Computational intelligence
- Computational complexity theory
- Information theory

- Philosophy
- Psychology and neurobiology
- Bayesian methods
- Mathematics
 - Probability
 - Statistics
 - Linear algebra
 - Functional analysis
 - Numerical analysis
- ...

2. Types of learning

1. Supervised learning (SL)

- **predictive models**
- **applications**
 - intrusion detection, data rectification for process control, image (pattern) recognition, predictions
- completely labeled training data
 - A trainer submits the input/output exemplary patterns and the learner has to adjust the parameters of the system autonomously, so that it can yield the correct output patterns when faced with a new input pattern.
 - A set of training examples $(x_i, f(x_i))$, where f is the target function to be learned, is provided to the learner and the aim is to determine (an approximation of) f by some adaptive algorithm.
- issues in supervised learning
 - **overfitting** (learn by heart)
 - The model learns very well the training data (it has high performance on the training data set), but it does not generalize well on unseen data (low performance on a testing data set)
 - **underfitting**
 - the model is too simple, it cannot capture the structure of the data
 - Noisy training data (errors in data, outliers, irrelevant data)
- 2 important types of supervised learning
 - **Inductive learning**
 - determine a hypothesis h such that $h(x_i) \approx f(x_i)$
 - how to compare two hypotheses approximately close to f ?
 - *inductive bias*
 - **Analogical learning**
 - Identifying analogies between an experienced problem instance and a new problem
 - E.g. case-based reasoning (CBR)

2. Unsupervised learning (UL)

- **descriptive models**
- completely unlabeled training data

- in absence of trainers, the desired output for a given input instance is not known, and consequently, the learner has to adapt its parameters autonomously.

3. Reinforcement learning (RL)

- learning by **interaction with the environment**
- an autonomous agent learns to perform an optimal sequence of actions to reach a goal
- **applications**
 - game playing, robotics and control

4. Semi-supervised learning

- falls between SL and UL
- the learner is provided with a small amount of labeled data and a large amount of unlabeled data
- graph-based, kernel-based, generative, pseudo-labeling methods, aso.

Inductive logic programming (ILP)

- subfield of ML which uses first-order logic to represent hypotheses and data
- specifically targets problems involving structured data and background knowledge
- tackles a wide variety of problems in machine learning, including classification, regression, clustering, and reinforcement learning.
- **applications**
 - in bio- and chemo-informatics, natural language processing, and web mining.

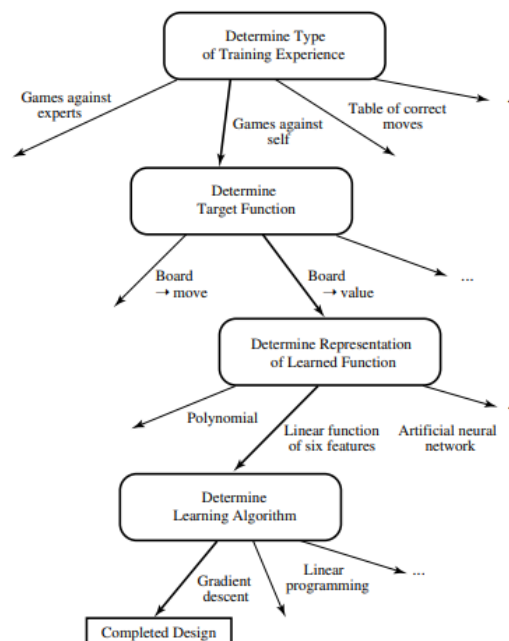
3. The learning problem

- **specification** of the learning task
 - T, P, E
 - improve over task T
 - with respect to performance measure P
 - based on experience E
 - e.g, learning to play a board game [\[1\]](#)
 - T : Play checkers
 - P : % of games won in world tournament
 - E : opportunity to play against self
 - what experience?
 - direct or indirect?
 - teacher or not?
 - is training experience representative to the performance goal?
- **design** choices
 - what exactly should be learned?
 - choose the **target function**
 - how shall the learned function be represented?
 - what specific algorithm to learn the target function?

Choose the Target Function

- $ChooseMove : Board \rightarrow Move$??
 - $V : Board \rightarrow \mathbb{R}$??
 - ...
- components of a learning system
 - the **performance system**
 - responsible with providing the output, using the learned target function
 - the **critic**
 - responsible with providing feedback to the learner (e.g., training examples in case of SL)
 - the **generalizer**
 - responsible with producing an output hypothesis that is the estimate of the target function
 - the **experiment generator**
 - mainly in RL scenarios
 - takes as input the current hypothesis (currently learned function) and outputs a new problem for the performance system to explore

Design Choices




[1]

- **Learning is most reliable when the training examples follow a distribution similar to that of future test examples**

- e.g., in game playing, the learner might never encounter certain crucial board states that are very likely to be played by the human champion.
- representative training data

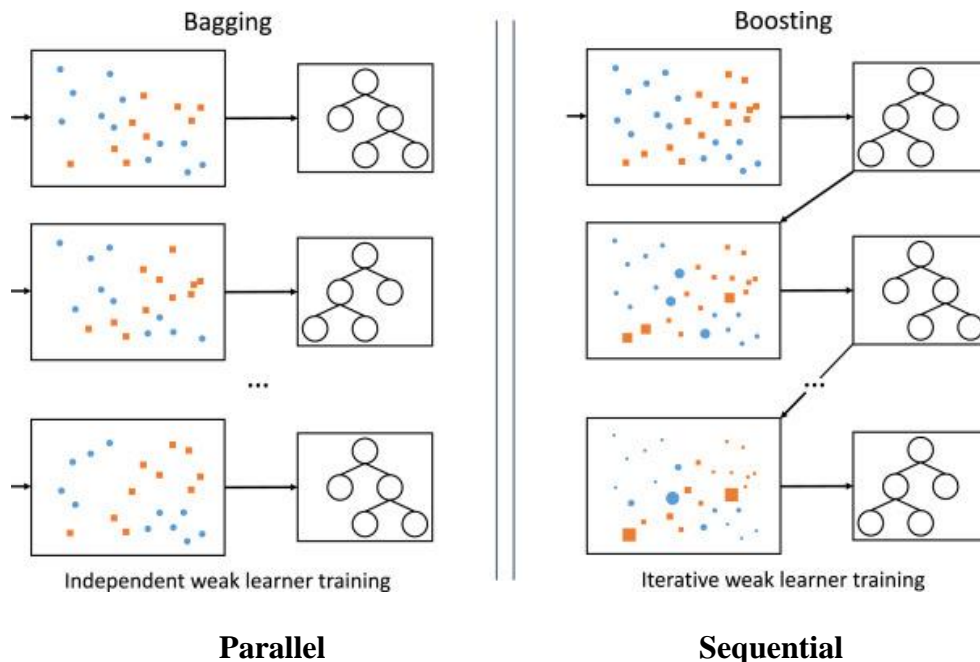
4. Ensemble learning

- meta-learning algorithms
- to combine multiple ML algorithms to obtain better predictive performance
- multiple models (often called “**weak learners**”) are trained to solve the same problem and combined to get better results.
 - the main assumption is that we can obtain more accurate and/or robust models when weak models are correctly combined.
- An ensemble learner outputs a single hypothesis which is not necessarily contained in the hypotheses space of the small models.
- types of [ensemble learning](#)
 - **Bagging**
 - ensemble meta-learning algorithm
 - improve stability and accuracy of ML models
 - reduces overfitting
 - **Bootstrapping AGG**regat**ING**
 - *bootstrapping* is random sampling **with replacement** (*an observation may be selected more than once*) from the available training data.
 - **Bagging** is performing bootstrapping many times and training an estimator for each bootstrapped dataset.
 - theoretical foundation
 - sampling with replacement and then building an ensemble reduces the variance of the ensemble learner without increasing the bias.
 - homogeneous weak learners
 - E.g. [Random Forests](#) (weak learner = decision tree)



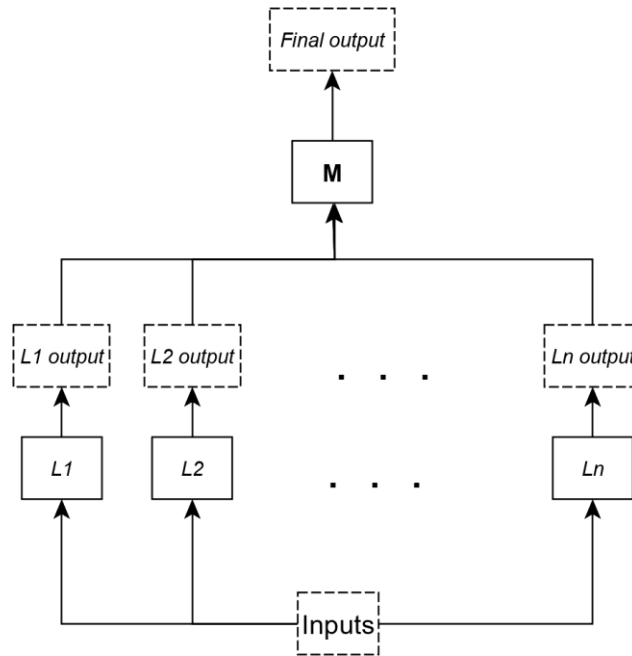
○ Boosting

- ensemble meta-learning algorithm
- use multiple homogeneous weak learners to obtain a single strong learner
 - a good weak learner is one that is highly biased (*unstable*): e.g., decision tree
- used to **reduce errors** in predictive data analysis
- can a set of weak learners slightly correlated with the true classification to create a single strong learner?
- seems to be better than bagging, but has the tendency to overfit
- **algorithms**: Adaboost, EpsilonBoost, Gradient boosting (XGBoost)



○ Stacking

- ensemble meta-learning algorithm
- heterogeneous weak learners
- combine the weak learners by training a metamodel



- ensemble of DL models

5. Probability, statistics and information theory

- **Probability & statistics**

- modelling processes with uncertainty

- **Statistics – inductive**

- we **observe** something that has happened and
- we try to figure out what **underlying process** would explain those observations

- **Probability – deductive**

- we consider some **underlying process** which has some randomness/uncertainty modelled by random variables and
- we try to figure out **what** happens

- random variables


- a variable whose possible values are outcomes of a random phenomenon
 - discrete (e.g. throwing a die)
 - continuous (e.g. the age of death)
- ML
 - features
 - outcome of the learning process

- **probability distributions** [2]

- describes probabilities of values a random variable can take
 - *discrete* (e.g. throwing a die)
 - *continuous* (e.g. the age of death)
 - continuous values (real numbers)


- **normal (Gaussian)**

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$



Source: https://commons.wikimedia.org/wiki/File:Normal_Distribution_PDF.svg

- 68.26% of the data fall within 1 SD of the mean
- 95.44% of the data fall within 2 SDs of the mean
 - 95% of the data fall within 1.96 SD of the mean




- 99.7% of the data fall within 3 SD of the mean
 - this rule enables to
 - check for outliers
 - determine the normality of any distribution
- In **Machine Learning**, data satisfying **Normal Distribution** is beneficial for model building.
 - models like LDA, Gaussian NBC, Logistic regression, Linear Regression, etc, use the assumption of normally distributed data.
 - sigmoid functions work most naturally with normally distributed data
 - natural phenomena (e.g., financial data, forecasting data) follow a *log-normal distribution*
 - continuous probability distribution of a random variable whose *logarithm* is *normally distributed*.
 - many processes follow **normality**
 - many measurement errors in an experiment
 - the position of a particle that experiences diffusion
 - ...
 - convert the data into a normal distribution - transformation techniques (statistical software packages)
 - e.g., **quantile transforms**

- will map a variable's probability distribution to another probability distribution
 - **QuantileTransformer** class from *scikit-learn* Python library
- **Normality** assumption for the ML models
 - it is not mandatory that data should always follow normality
 - ML models may work very well in the case of non-normally distributed data
 - **Decision trees, XGBoost**, don't assume any normality and work on raw data.
 - **Linear regression** is statistically effective if only the model errors are Gaussian, not exactly the entire dataset.
- **Statistical tests** for checking data normality
 - **Shapiro-Wilk** test, **Kolmogorov-Smirnov** (K-S) test, ..
 - statistical software

!!! Explore the data and check for the underlying distributions for each variable/feature before going to fit the ML model.

- **moments**
 - [mean, variance, standard deviation](#)
 - **mean** – expected value
 - $\mu, E(X)$
 - **variance** – expected squared distance from mean
 - $\sigma^2 = E((X - E(X))^2)$
 - **standard deviation**
 - σ
 - measures spread of distribution
 - how far a set of numbers is spread out
 - low σ - close to the mean
 - **Central limit theorem**
 - the sum of a large number of random samples is normally distributed
 - the sum of N random variables with mean m and standard deviation s can be approximated by a normal distribution with
 - mean $\mu = N \cdot m$
 - standard deviation $\sigma = s \cdot \sqrt{N}$
 - if one takes sufficiently large samples from a population, the samples' means will be *normally distributed*, even if the population isn't normally distributed
 - application: **confidence intervals**
- comparing the performance of ML classifiers/regressors
 - **statistical significance tests**
 - T-test, Wilcoxon test, etc
- **correlation** – statistical relationship between two random variables
 - intuition
 - two random variables X and Y are correlated, when:
 - when X is high, Y is likely to be high
 - when X is low, Y is likely to be low
 - **correlation coefficients** – measure the degree to which the variables are monotonically related

- **Pearson**
 - is used for measuring the degree of linear relationship between two features
- **Spearman**
 - describes the strength and direction of the monotonic relationship between two variables (even if their relationship is not linear)
- correlation coefficient formulas are used to find how strong a relationship is between data
- return a value between -1 and 1, where:
 - 1 indicates a strong positive relationship.
 - -1 indicates a strong negative relationship.
 - 0 indicates no relationship at all.



Source: <https://www.statisticshowto.com/probability-and-statistics/correlation-coefficient-formula/>

- in supervised learning
 - features **well correlated** with the target output
 - **independent** features
- [statistics](#) [2]: statistical independence, confidence intervals
 - two random variables are *independent* if they convey no information about each other

$$P(A|B) = P(A)$$
 - Pearson's χ^2 (Chi-squared) test for independence
 - **independence** \Rightarrow uncorrelation, but uncorrelation does not imply independence
 - choosing relevant and independent features is key to ML
 - **confidence interval** (CI)
 - a statistical measure of the reliability/consistency of an estimate
 - is an interval estimate of a population parameter (e.g. the mean)
 - an interval that frequently includes the parameter
 - 95% CI of the mean μ of a population
 - $[\mu - \alpha, \mu + \alpha]$ or $\mu \pm \alpha$
 - α - *confidence value*
 - $\alpha = 1.96 * \sigma / \sqrt{N}$
 - N – sample size
 - σ - population standard deviation
 - comparing the performance of classifiers

- conditional probability (Bayes theorem)

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

- **Bayesian learning**
- **entropy**
 - a measure from information theory
 - measures the impurity of a set

[SLIDES]

- [Introduction to ML](#) [1]

[READING]

- [Introduction to ML](#) (T. Mitchell) [1]
- [Introduction to ML and DL](#) (Zhang et al.) [2]
- [ML preliminaries](#) (N. Nillson) [4]
- [ML Basics](#) (Goodfellow et al.) [3]
- [Probability and information theory](#) (Goodfellow et al.) [3]
- [Evaluating hypotheses](#) (T. Mitchell) [1]

Bibliography

[1] Mitchell, T., *Machine Learning*, McGraw Hill, 1997 (available at www.cs.ubbcluj.ro/~gabis/ml/ml-books)

[2] Aston Zhang, Zachary C. Lipton, Mu Li, and Alexander J. Smola, *Dive into Deep Learning*, 2020 (<http://d2l.ai/>)

[3] Ian Goodfellow, Yoshua Bengio, Aaron Courville, *Deep Learning*, MIT Press, 2016 (online edition at <http://www.deeplearningbook.org/>)

[4] Nillson, N., *Introduction to Machine Learning*, Stanford University, 1996 (available at www.cs.ubbcluj.ro/~gabis/ml/ml-books)